

## ▾ Getting started with neural networks: Classification and regression

### Classifying movie reviews: A binary classification example

#### The IMDB dataset

#### Loading the IMDB dataset

```
from tensorflow.keras.datasets import imdb
(train_data, train_labels), (test_data, test_labels) = imdb.load_data(
    num_words=10000)

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb.npz
17464789/17464789 [=====] - 0s 0us/step
```

```
train_data[0]
```

```
104,
88,
4,
381,
15,
297,
98,
32,
2071,
56,
26,
141,
6,
194,
7486,
18,
4,
226,
22,
21,
134,
476,
26,
480,
5,
144,
30,
5535,
18,
51,
36,
28,
224,
92,
25,
104,
4,
226,
65,
16,
38,
1334,
88,
12,
16,
283,
5,
16,
4472,
113,
103,
32,
15,
16,
5345,
19,
178,
32]
```

```
train_labels[0]
```

1

```
max([max(sequence) for sequence in train_data])
```

9999

### Decoding reviews back to text

```
word_index = imdb.get_word_index()
reverse_word_index = dict(
    [(value, key) for (key, value) in word_index.items()])
decoded_review = " ".join(
    [reverse_word_index.get(i - 3, "?") for i in train_data[0]])
```

Downloading data from [https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb\\_word\\_index.json](https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb_word_index.json)  
 1641221/1641221 [=====] - 0s 0us/step

## ▾ Preparing the data

### Encoding the integer sequences via multi-hot encoding

```
import numpy as np
def vectorize_sequences(sequences, dimension=10000):
    results = np.zeros((len(sequences), dimension))
    for i, sequence in enumerate(sequences):
        for j in sequence:
            results[i, j] = 1.
    return results
x_train = vectorize_sequences(train_data)
x_test = vectorize_sequences(test_data)
```

```
x_train[0]

array([0., 1., 1., ..., 0., 0., 0.])
```

```
y_train = np.asarray(train_labels).astype("float32")
y_test = np.asarray(test_labels).astype("float32")
```

## ▾ Building your model

### Model definition and compilation

```
from tensorflow import keras
from tensorflow.keras import layers

model = keras.Sequential([
    layers.Dense(32, activation="relu"),
    layers.Dense(32, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])

model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
```

### ▾ Validating the data set

```
x_val = x_train[:10000]
partial_x_train = x_train[10000:]
y_val = y_train[:10000]
partial_y_train = y_train[10000:]
```

### ▾ Training model

```

history = model.fit(partial_x_train,
                    partial_y_train,
                    epochs=25,
                    batch_size=512,
                    validation_data=(x_val, y_val))

Epoch 1/25
30/30 [=====] - 7s 81ms/step - loss: 0.5177 - accuracy: 0.7625 - val_loss: 0.3667 - val_accuracy: 0.8706
Epoch 2/25
30/30 [=====] - 1s 35ms/step - loss: 0.3009 - accuracy: 0.8939 - val_loss: 0.3735 - val_accuracy: 0.8386
Epoch 3/25
30/30 [=====] - 1s 23ms/step - loss: 0.2266 - accuracy: 0.9187 - val_loss: 0.2903 - val_accuracy: 0.8832
Epoch 4/25
30/30 [=====] - 1s 24ms/step - loss: 0.1901 - accuracy: 0.9313 - val_loss: 0.2874 - val_accuracy: 0.8851
Epoch 5/25
30/30 [=====] - 1s 22ms/step - loss: 0.1563 - accuracy: 0.9461 - val_loss: 0.3043 - val_accuracy: 0.8804
Epoch 6/25
30/30 [=====] - 1s 23ms/step - loss: 0.1307 - accuracy: 0.9558 - val_loss: 0.3344 - val_accuracy: 0.8729
Epoch 7/25
30/30 [=====] - 1s 25ms/step - loss: 0.1126 - accuracy: 0.9610 - val_loss: 0.3236 - val_accuracy: 0.8784
Epoch 8/25
30/30 [=====] - 1s 25ms/step - loss: 0.0965 - accuracy: 0.9687 - val_loss: 0.3335 - val_accuracy: 0.8796
Epoch 9/25
30/30 [=====] - 1s 22ms/step - loss: 0.0861 - accuracy: 0.9708 - val_loss: 0.3706 - val_accuracy: 0.8737
Epoch 10/25
30/30 [=====] - 1s 22ms/step - loss: 0.0718 - accuracy: 0.9772 - val_loss: 0.3782 - val_accuracy: 0.8778
Epoch 11/25
30/30 [=====] - 1s 22ms/step - loss: 0.0608 - accuracy: 0.9810 - val_loss: 0.4041 - val_accuracy: 0.8755
Epoch 12/25
30/30 [=====] - 1s 25ms/step - loss: 0.0581 - accuracy: 0.9829 - val_loss: 0.4221 - val_accuracy: 0.8755
Epoch 13/25
30/30 [=====] - 1s 24ms/step - loss: 0.0363 - accuracy: 0.9915 - val_loss: 0.6675 - val_accuracy: 0.8395
Epoch 14/25
30/30 [=====] - 1s 25ms/step - loss: 0.0318 - accuracy: 0.9921 - val_loss: 0.4793 - val_accuracy: 0.8729
Epoch 15/25
30/30 [=====] - 1s 21ms/step - loss: 0.0272 - accuracy: 0.9941 - val_loss: 0.6724 - val_accuracy: 0.8448
Epoch 16/25
30/30 [=====] - 1s 22ms/step - loss: 0.0270 - accuracy: 0.9932 - val_loss: 0.5387 - val_accuracy: 0.8718
Epoch 17/25
30/30 [=====] - 1s 38ms/step - loss: 0.0305 - accuracy: 0.9915 - val_loss: 0.5618 - val_accuracy: 0.8703
Epoch 18/25
30/30 [=====] - 1s 29ms/step - loss: 0.0108 - accuracy: 0.9994 - val_loss: 0.5988 - val_accuracy: 0.8698
Epoch 19/25
30/30 [=====] - 1s 26ms/step - loss: 0.0247 - accuracy: 0.9931 - val_loss: 0.6146 - val_accuracy: 0.8696
Epoch 20/25
30/30 [=====] - 1s 25ms/step - loss: 0.0206 - accuracy: 0.9943 - val_loss: 0.6350 - val_accuracy: 0.8692
Epoch 21/25
30/30 [=====] - 1s 21ms/step - loss: 0.0059 - accuracy: 0.9997 - val_loss: 0.8964 - val_accuracy: 0.8408
Epoch 22/25
30/30 [=====] - 1s 22ms/step - loss: 0.0231 - accuracy: 0.9929 - val_loss: 0.6661 - val_accuracy: 0.8693
Epoch 23/25
30/30 [=====] - 1s 22ms/step - loss: 0.0165 - accuracy: 0.9951 - val_loss: 0.6982 - val_accuracy: 0.8683
Epoch 24/25
30/30 [=====] - 1s 21ms/step - loss: 0.0039 - accuracy: 0.9998 - val_loss: 0.7090 - val_accuracy: 0.8683
Epoch 25/25
30/30 [=====] - 1s 20ms/step - loss: 0.0173 - accuracy: 0.9945 - val_loss: 0.7303 - val_accuracy: 0.8682

history_dict = history.history
history_dict.keys()

dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])


```

## ▼ Training and validation loss

```


import matplotlib.pyplot as plt
history_dict = history.history
loss_values = history_dict["loss"]
val_loss_values = history_dict["val_loss"]
epochs = range(1, len(loss_values) + 1)
plt.plot(epochs, loss_values, "bo", label="Training loss")
plt.plot(epochs, val_loss_values, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()

```



### ▼ Training and validation accuracy

```
plt.clf()
acc = history_dict["accuracy"]
val_acc = history_dict["val_accuracy"]
plt.plot(epochs, acc, "bo", label="Training acc")
plt.plot(epochs, val_acc, "b", label="Validation acc")
plt.title("Training and validation accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
```



### ▼ Retraining a model from scratch

```
model = keras.Sequential([
    layers.Dense(32, activation="relu"),
    layers.Dense(32, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.fit(x_train, y_train, epochs=10, batch_size=512)
```

```
results = model.evaluate(x_test, y_test)
```

```
Epoch 1/10
49/49 [=====] - 2s 19ms/step - loss: 0.4425 - accuracy: 0.8115
Epoch 2/10
49/49 [=====] - 1s 16ms/step - loss: 0.2547 - accuracy: 0.9066
Epoch 3/10
49/49 [=====] - 1s 18ms/step - loss: 0.2070 - accuracy: 0.9219
Epoch 4/10
49/49 [=====] - 1s 14ms/step - loss: 0.1760 - accuracy: 0.9350
Epoch 5/10
49/49 [=====] - 1s 14ms/step - loss: 0.1546 - accuracy: 0.9437
Epoch 6/10
49/49 [=====] - 1s 14ms/step - loss: 0.1344 - accuracy: 0.9518
Epoch 7/10
49/49 [=====] - 1s 14ms/step - loss: 0.1189 - accuracy: 0.9572
Epoch 8/10
49/49 [=====] - 1s 13ms/step - loss: 0.1016 - accuracy: 0.9640
Epoch 9/10
49/49 [=====] - 1s 14ms/step - loss: 0.0885 - accuracy: 0.9696
Epoch 10/10
49/49 [=====] - 1s 12ms/step - loss: 0.0753 - accuracy: 0.9750
782/782 [=====] - 3s 3ms/step - loss: 0.4232 - accuracy: 0.8688
```

```
results
```

```
[0.42323464155197144, 0.8688399791717529]
```